

MINISTRY OF EDUCATION AND TRAINING
EASTERN INTERNATIONAL UNIVERSITY



MIS 443
BUSINESS DATA MANAGEMENT

Final Project
Case Study 4 - Data Bank

Lecturers: *Mr. Dang Thai Doan*
Ms. Huynh Tuyet Ngan

GROUP 3A	
Name	IRN
Nguyễn Thành Đạt	2132300562
Lê Nguyễn Tâm Như	2132300065
Nguyễn Hoàng Vinh	2132300522

Quarter 4/2024-2025

Date Submission: 13/8/2025

Table Of Content

1. Objective.....	3
2. Project Overview: Data Bank.....	3
2.1. Introduction.....	3
2.2. Key Data Elements.....	3
3. Deliverables.....	3
A. Customer Nodes Exploration.....	3
B. Customer Transactions.....	8
C. Data Allocation Challenge.....	13
4. Final Summary.....	18
4.1. Customer Nodes.....	18
4.2. Customer Transactions.....	18
4.3. Data Allocation Challenge.....	18

1. Objective

The objective of this report is to analyze the Data Bank case study by applying SQL queries to address key business questions. The analysis involves understanding customer behavior, transaction patterns, and regional performance using real-world data. The aim is to extract actionable insights that align with business objectives, which can inform decisions about customer segmentation, regional marketing strategies, and operational optimizations. This report highlights the SQL skills employed to develop queries and derive meaningful conclusions from the data, focusing on enhancing analytical thinking and problem-solving abilities in a business context.

2. Project Overview: Data Bank

2.1. Introduction

The "Data Bank" project simulates a digital-only bank integrated with a secure distributed data storage platform. The bank allocates data storage limits based on the amount of money customers hold in their accounts. The aim of the project is to analyze various datasets related to customer transactions and data allocations, and to provide insights into Data Bank's operations. This analysis will help in better understanding the customer base, forecasting future storage needs, and addressing core business questions for improved strategic planning.

In this project, we utilize SQL queries to explore, aggregate, and analyze data from three main tables: regions, customer_nodes, and customer_transactions. The goal is to uncover patterns, summarize important metrics, and answer business-critical questions that will assist in Data Bank's growth and operational efficiency.

2.2. Key Data Elements

- Regions - Information about different geographical regions where Data Bank operates.
- Customer Nodes - Details on customer allocation to specific nodes (data storage locations) along with their active dates.
- Customer Transactions - Records of customer financial activities, including deposits, withdrawals, and purchases made using their Data Bank debit card.

3. Deliverables

A. Customer Nodes Exploration

Question 1: How many unique nodes are there on the Data Bank system?

SQL Query:

```

5  ▼ select
6      count(distinct node_id) as total_unique_nodes
7  from
8      data_bank.customer_nodes;

```

SQL output:

	total_unique_nodes 
1	5

Insights:

The Data Bank system has 5 unique nodes, indicating a distributed infrastructure that enhances data security and redundancy. This setup ensures efficient storage management, reduces the risk of data loss, and supports scalability as the customer base grows.

Question 2 What is the number of nodes per region?

SQL Query:

```

11 ▼ select
12     rg.region_name, count(distinct nd.node_id) as total_nodes
13 from
14     data_bank.customer_nodes as nd
15 inner join
16     data_bank.regions as rg on rg.region_id = nd.region_id
17 group by
18     rg.region_name
19 order by
20     rg.region_name;

```

SQL output:

	region_name character varying (9) 🔒	total_nodes bigint 🔒
1	Africa	5
2	America	5
3	Asia	5
4	Australia	5
5	Europe	5

Insights:

Each region in the Data Bank system contains exactly 5 unique nodes. This uniform distribution across all regions suggests that Data Bank aims for equal infrastructural capacity globally. It reflects a strategy to ensure consistent data storage and security, reducing regional disparities and providing balanced access to their services across different regions.

Question 3: How many customers are allocated to each region?

SQL Query:

```
23  ✓ select
24      rg.region_name, count(distinct nd.customer_id) as total_customers
25  from
26      data_bank.customer_nodes as nd
27  inner join
28      data_bank.regions as rg on rg.region_id = nd.region_id
29  group by
30      rg.region_name
31  order by
32      total_customers desc;
```

SQL output:

	region_name character varying (9) 🔒	total_customers bigint 🔒
1	Australia	110
2	America	105
3	Africa	102
4	Asia	95
5	Europe	88

Insights:

The number of customers allocated to each region varies, with Australia having the highest at 110 customers and Europe having the lowest at 88 customers. This distribution suggests that customer density is higher in certain regions, likely due to factors such as regional popularity or market penetration. Regions with fewer customers may require targeted strategies to increase customer acquisition and engagement.

Question 4: How many days on average are customers reallocated to a different node?

SQL Query:

```
38 WITH active_days_in_nodes AS
39     (SELECT
40         nd.customer_id,
41         nd.node_id,
42         SUM(nd.end_date - nd.start_date) AS active_days_in_nodes
43     FROM
44         data_bank.customer_nodes AS nd
45     WHERE
46         end_date != '9999-12-31'
47     GROUP BY
48         nd.customer_id, nd.node_id
49     ORDER BY
50         customer_id, node_id)
51 SELECT round(AVG(adin.active_days_in_nodes)) AS active_days_in_nodes
52 FROM active_days_in_nodes AS adin
```

SQL output:

	active_days_in_nodes numeric
1	24

Insights:

On average, customers are reallocated to a different node every 24 days. This reflects how often Data Bank reassigned customers to new storage nodes, possibly due to their security protocols or data management strategies. Regular node reallocation ensures that data is distributed evenly and reduces the risk of data breaches or system vulnerabilities, providing customers with enhanced security and reliable service.

Question 5: What is the median, 80th and 95th percentile for this same reallocation days metric for each region?

SQL Query:

```

56 SELECT
57     rg.region_name,
58     PERCENTILE_CONT(0.5) WITHIN GROUP (ORDER BY nd.end_date - nd.start_date) AS median_days,
59     PERCENTILE_CONT(0.8) WITHIN GROUP (ORDER BY nd.end_date - nd.start_date) AS p80_days,
60     PERCENTILE_CONT(0.95) WITHIN GROUP (ORDER BY nd.end_date - nd.start_date) AS p95_days
61 FROM
62     data_bank.customer_nodes as nd
63 INNER JOIN
64     data_bank.regions as rg on rg.region_id = nd.region_id
65 WHERE
66     end_date != '9999-12-31'
67 GROUP BY
68     rg.region_name
69 ORDER BY
70     rg.region_name;

```

SQL output:

	region_name character varying (9)	median_days double precision	p80_days double precision	p95_days double precision
1	Africa	15	24	28
2	America	15	23	28
3	Asia	15	23	28
4	Australia	15	23	28
5	Europe	15	24	28

Insights:

The median, 80th, and 95th percentiles for reallocation days vary slightly across regions. The median is consistently 15 days for all regions, while the 80th percentile ranges from 23 to 24 days, and the 95th percentile remains 28 days. This indicates that while most customers are reallocated within 15 days, some regions experience slightly longer reallocation periods, especially for a small portion of customers at the 80th percentile. These variations suggest that Data Bank's reallocation strategy may differ slightly by region, likely influenced by factors such as data load or operational differences.

B. Customer Transactions

Question 1: What is the unique count and total amount for each transaction type?

SQL Query:

```
76  SELECT
77      txn_type,
78      COUNT(customer_id) as total_transactions,
79      sum(txn_amount) as total_amount
80  FROM
81      data_bank.customer_transactions
82  GROUP BY
83      txn_type
84  ORDER BY
85      txn_type;
```

SQL output:

	txn_type character varying (10) 🔒	total_transactions bigint 🔒	total_amount bigint 🔒
1	deposit	2671	1359168
2	purchase	1617	806537
3	withdrawal	1580	793003

Insights:

The analysis shows that deposit transactions are the most frequent, with 2,671 occurrences and a total of 1,359,168. Purchase transactions follow with 1,617 occurrences and 806,537, while withdrawals have 1,580 occurrences totaling 793,003. This indicates that deposits are the primary activity in the system, with purchases and withdrawals representing balanced

financial behavior. The high deposit total underscores the focus on customer savings and account growth.

Question 2: What is the average total historical deposit counts and amounts for all customers?

SQL Query:

```
89  SELECT
90      ROUND(AVG(deposit_count)) AS avg_deposit_times,
91      ROUND(AVG(deposit_amount)) AS avg_deposit_amount
92  FROM (
93      SELECT
94          customer_id,
95          COUNT(*) AS deposit_count,
96          AVG(txn_amount) AS deposit_amount
97  FROM data_bank.customer_transactions
98  WHERE txn_type = 'deposit'
99  GROUP BY customer_id
100 ) AS deposit;
```

SQL output:

	avg_deposit_times numeric	avg_deposit_amount numeric
1	5	509

Insights:

On average, each customer has made 5 deposits historically, with an average deposit amount of 509. This suggests that customers tend to make a moderate number of deposits over time, with relatively consistent transaction sizes. The deposit amount indicates typical individual deposit values, reflecting steady but not excessively large contributions into customer accounts.

Question 3: For each month - how many Data Bank customers make more than 1 deposit and either 1 purchase or 1 withdrawal in a single month?

SQL Query:

```
106 WITH customer_dep_pur_with_count AS (SELECT
107     customer_id,
108     EXTRACT(MONTH FROM txn_date) AS month,
109     SUM(CASE WHEN txn_type = 'deposit' then 1 else 0 end) as total_deposit_count,
110     SUM(CASE WHEN txn_type = 'purchase' then 1 else 0 end) as total_purchase_count,
111     SUM(CASE WHEN txn_type = 'withdrawal' then 1 else 0 end) as total_withdrawal_count
112 FROM
113     data_bank.customer_transactions
114 GROUP BY
115     customer_id, month
116 ORDER BY
117     customer_id, month)
118 SELECT
119     MONTH, COUNT(DISTINCT customer_id) as total_customers
120 FROM
121     customer_dep_pur_with_count
122 WHERE
123     total_deposit_count > 1 AND (total_purchase_count >= 1 OR total_withdrawal_count >=1)
124 GROUP BY month
125 ORDER BY month;
```

SQL output:

	month numeric 🔒	total_customers bigint 🔒
1	1	168
2	2	181
3	3	192
4	4	70

Insights:

The results show how many customers made more than 1 deposit and at least 1 purchase or withdrawal in each month. In the first month (January), 168 customers met this criterion, while 181 customers did so in February and 192 customers in March. The data indicates a growing trend of active engagement among customers, with a noticeable increase in transaction activity, likely reflecting more frequent use of their accounts for deposits and other transactions as the months progress.

Question 4: What is the closing balance for each customer at the end of the month?

SQL Query:

```

128 SELECT
129     customer_id,
130     txn_month,
131     SUM(net_change) OVER (
132         PARTITION BY customer_id
133         ORDER BY txn_month
134     ) AS closing_balance
135 FROM (
136     SELECT
137         customer_id,
138         DATE_TRUNC('month', txn_date) AS txn_month,
139         SUM(
140             CASE
141                 WHEN txn_type = 'deposit' THEN txn_amount
142                 WHEN txn_type IN ('withdrawal', 'purchase') THEN -txn_amount
143                 ELSE 0
144             END
145         ) AS net_change
146     FROM data_bank.customer_transactions
147     GROUP BY customer_id, DATE_TRUNC('month', txn_date)
148 ) AS monthly_change
149 ORDER BY customer_id, txn_month;

```

SQL output:

	customer_id integer	txn_month timestamp with time zone	closing_balance numeric
1	1	2020-01-01 00:00:00+07	312
2	1	2020-03-01 00:00:00+07	-640
3	2	2020-01-01 00:00:00+07	549
4	2	2020-03-01 00:00:00+07	610
5	3	2020-01-01 00:00:00+07	144
6	3	2020-02-01 00:00:00+07	-821
7	3	2020-03-01 00:00:00+07	-1222
8	3	2020-04-01 00:00:00+07	-729
9	4	2020-01-01 00:00:00+07	848
10	4	2020-03-01 00:00:00+07	655
11	5	2020-01-01 00:00:00+07	954
12	5	2020-03-01 00:00:00+07	-1923
13	5	2020-04-01 00:00:00+07	-2413
14	6	2020-01-01 00:00:00+07	733
15	6	2020-02-01 00:00:00+07	-52

Insights:

The SQL query returns each customer's closing balance at the end of every month, with positive values indicating net deposits and negative values indicating net withdrawals or spending. From the output, we can observe varied financial patterns: some customers maintain positive balances across months, while others experience significant declines, sometimes turning negative. These fluctuations highlight differences in spending habits, saving behavior, and potential financial risk, which could inform targeted financial advice or interventions.

Question 5: What is the percentage of customers who increase their closing balance by more than 5%?

SQL Query:

```
151 WITH monthly_closing_balance AS (  
152     SELECT  
153         customer_id,  
154         EXTRACT(MONTH FROM txn_date) AS month,  
155         SUM(CASE  
156             WHEN txn_type = 'deposit' THEN txn_amount  
157             WHEN txn_type IN ('withdrawal', 'purchase') THEN -txn_amount  
158             ELSE 0  
159         END) AS balance  
160     FROM data_bank.customer_transactions  
161     GROUP BY customer_id, month  
162     ORDER BY customer_id, month  
163 ),  
  
164 |  
165 balance_with_previous AS (  
166     SELECT  
167         customer_id,  
168         month,  
169         balance,  
170         LAG(balance) OVER (PARTITION BY customer_id ORDER BY month) AS prev_balance  
171     FROM monthly_closing_balance  
172 ),  
173 increase_over_5_percent as (SELECT *  
174     FROM balance_with_previous  
175     WHERE prev_balance is not null AND balance > prev_balance and ((balance/prev_balance) > 1.05 )  
176 ),  
177 final_percentage as (SELECT  
178     COUNT(DISTINCT customer_id) * 100.0 /  
179     (SELECT COUNT(DISTINCT customer_id) FROM data_bank.customer_transactions) AS percent_increased  
180     FROM increase_over_5_percent)  
181  
182 SELECT round(percent_increased,2) || '%' as increase_over_5_percent FROM final_percentage;
```

SQL output:

	increase_over_5_percent text
1	13.20%

Insights:

Approximately 13.20% of customers have increased their closing balance by more than 5% in a given month. This indicates that a portion of customers are seeing steady growth in their account balances, which could reflect successful savings or positive financial activity.

Monitoring this metric is valuable for Data Bank to assess customer engagement and identify strategies to help more customers achieve similar growth.

C. Data Allocation Challenge

Question: To test out a few different hypotheses - the Data Bank team wants to run an experiment where different groups of customers would be allocated data using 3 different options:

- ***Option 1: data is allocated based off the amount of money at the end of the previous month***
- ***Option 2: data is allocated on the average amount of money kept in the account in the previous 30 days***
- ***Option 3: data is updated real-time***

For this multi-part challenge question - you have been requested to generate the following data elements to help the Data Bank team estimate how much data will need to be provisioned for each option:

- ***running customer balance column that includes the impact each transaction***
- ***customer balance at the end of each month***
- ***minimum, average and maximum values of the running balance for each customer***

Using all of the data available - how much data would have been required for each option on a monthly basis?

a. Running customer balance column that includes the impact each transaction

SQL Query:

```

1  -- running customer balance column that includes the impact each transaction
2  v SELECT customer_id,
3         txn_date,
4         txn_type,
5         txn_amount,
6         SUM(CASE
7             WHEN txn_type = 'deposit' THEN txn_amount
8             WHEN txn_type IN ('withdrawal', 'purchase') THEN - txn_amount
9             ELSE 0
10          END) OVER (PARTITION BY customer_id ORDER BY txn_date
11                    ROWS BETWEEN UNBOUNDED PRECEDING AND CURRENT ROW) AS running_balance
12 FROM data_bank.customer_transactions;

```

SQL output:

	customer_id [PK] integer	txn_date [PK] date	txn_type [PK] character varying (10)	txn_amount [PK] integer	running_balance bigint
1	1	2020-01-02	deposit	312	312
2	1	2020-03-05	purchase	612	-300
3	1	2020-03-17	deposit	324	24
4	1	2020-03-19	purchase	664	-640
5	2	2020-01-03	deposit	549	549
6	2	2020-03-24	deposit	61	610
7	3	2020-01-27	deposit	144	144
8	3	2020-02-22	purchase	965	-821
9	3	2020-03-05	withdrawal	213	-1034
10	3	2020-03-19	withdrawal	188	-1222
11	3	2020-04-12	deposit	493	-729
12	4	2020-01-07	deposit	458	458
13	4	2020-01-21	deposit	390	848
14	4	2020-03-25	purchase	193	655
15	5	2020-01-15	deposit	874	874

Insights:

This query calculates the running balance for each customer after each transaction. It reflects how each transaction (deposit, withdrawal, or purchase) impacts the customer's account balance over time. The results show that customers with positive balances, such as customer_id 1, indicate a healthy financial state, while customers with negative balances, such as customer_id 3 and customer_id 4, reflect scenarios where withdrawals or purchases have exceeded deposits. These variations highlight the importance of dynamic data allocation,

as customers with fluctuating or negative balances may require different storage capacities based on the volume and frequency of their transactions.

b. Customer balance at the end of each month

SQL Query:

```

13 -- customer balance at the end of each month
14 v SELECT
15     customer_id,
16     end_of_month,
17     SUM(net_change)
18     OVER (PARTITION BY customer_id
19           ORDER BY end_of_month) AS closing_balance
20 FROM (
21     SELECT
22         customer_id,
23         EXTRACT(MONTH from txn_date) AS end_of_month,
24         SUM(
25             CASE
26                 WHEN txn_type = 'deposit' THEN txn_amount
27                 WHEN txn_type IN ('withdrawal', 'purchase') THEN -txn_amount
28                 ELSE 0
29             END
30         ) AS net_change
31     FROM data_bank.customer_transactions
32     GROUP BY customer_id, end_of_month
33     ORDER BY customer_id, end_of_month
34 ) AS monthly_change
35 ORDER BY customer_id, end_of_month;

```

SQL output:

	customer_id integer	end_of_month numeric	closing_balance numeric
1	1	1	312
2	1	3	-640
3	2	1	549
4	2	3	610
5	3	1	144
6	3	2	-821
7	3	3	-1222

Insights:

This query calculates the closing balance for each customer at the end of each month. For example, customer 1 has a closing balance of 312 for January, while customer 3 shows a negative balance of -1222 for March. This reflects how the balance in each account fluctuates over time, indicating either an increase in savings or a decrease due to withdrawals or purchases. These fluctuations in monthly balances are essential for determining the data storage needs for each customer. Customers with higher balances may require more storage, while customers with negative or fluctuating balances might need less. By understanding these balance changes, Data Bank can more efficiently allocate data resources based on customer account activity and the level of financial fluctuations.

c. Minimum, average and maximum values of the running balance for each customer

SQL Query:

```
38  -- minimum, average and maximum values of the running balance for each customer
39
40  WITH running_balance AS (
41      SELECT customer_id,
42             txn_date,
43             txn_type,
44             txn_amount,
45             SUM(CASE
46                 WHEN txn_type = 'deposit' THEN txn_amount
47                 WHEN txn_type IN ('withdrawal', 'purchase') THEN -txn_amount
48                 ELSE 0
49             END) OVER (PARTITION BY customer_id ORDER BY txn_date) AS running_balance
50      FROM data_bank.customer_transactions
51  )
52  SELECT customer_id,
53         MIN(running_balance) AS min_balance,
54         AVG(running_balance) AS avg_balance,
55         MAX(running_balance) AS max_balance
56  FROM running_balance
57  GROUP BY customer_id
58  ORDER BY customer_id;
```


SQL output:

	customer_id integer	min_balance bigint	avg_balance numeric	max_balance bigint
1	1	-640	-151.000000000000000000	312
2	2	549	579.500000000000000000	610
3	3	-1222	-732.400000000000000000	144
4	4	458	653.666666666666666667	848
5	5	-2413	-135.454545454545454545	1780
6	6	-552	624.000000000000000000	2197
7	7	887	2268.6923076923076923	3539
8	8	-1029	173.700000000000000000	1363
9	9	-91	1021.700000000000000000	2030
10	10	-5090	-2229.833333333333333333	556
11	11	-2529	-1950.8235294117647059	60
12	12	-647	-14.500000000000000000	295
13	13	379	901.1538461538461538	1444
14	14	205	751.000000000000000000	989

Insights:

The query calculates the minimum, average, and maximum running balance for each customer, providing insights into how their account balance fluctuates over time. For example, customer 1 has a minimum balance of -640 and a maximum balance of 312, indicating a volatile account, while customer 2 has a stable positive average balance of 579.5 and a maximum balance of 610. The fluctuations in the balances reflect how customers manage their deposits, withdrawals, and purchases. Those with negative balances might experience more withdrawals than deposits, while customers with positive balances are likely maintaining savings. The minimum, average, and maximum values allow Data Bank to estimate storage needs, as customers with greater volatility or higher balances may require more data resources for tracking financial activity.

4. Final Summary

This report provides a comprehensive analysis of the Data Bank case study using SQL queries to answer key business questions, focusing on customer behavior, transaction patterns, and data allocation strategies.

4.1. Customer Nodes

Data Bank has an effective node distribution system across its regions, ensuring secure and efficient data storage. Customers are reallocated to different nodes approximately every 23-24 days, indicating a systematic and balanced approach to ensuring both data security and reliability.

4.2. Customer Transactions

Deposits are the most common type of transaction in the system, with customers regularly depositing funds. The average deposit per customer remains steady, and 13.20% of customers have seen their closing balance increase by more than 5% in a given month, indicating positive growth in their accounts.

4.3. Data Allocation Challenge

Three data allocation methods were explored:

- Option 1: Data allocation based on the end-of-month balance.
- Option 2: Data allocation based on the average balance over the past 30 days.
- Option 3: Real-time data updates for each transaction.

Each method offers unique advantages, with Option 2 providing a more dynamic allocation strategy that accounts for fluctuations in customer balances, making it a more flexible option for data management.